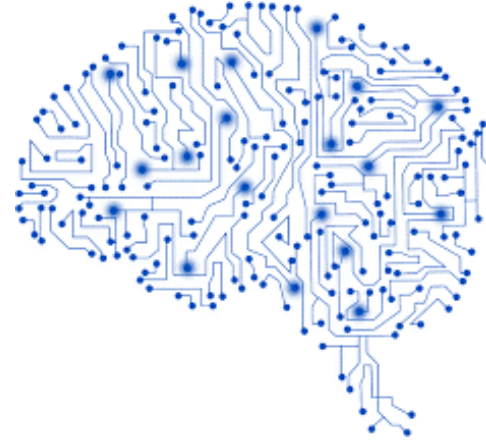




University of Minho
School of Engineering



Dados e Aprendizagem Automática

Unsupervised Learning:

K-means, K-medoids and DBSCAN

DAA @ MEI-1º/MiEI-4º – 1º Semestre

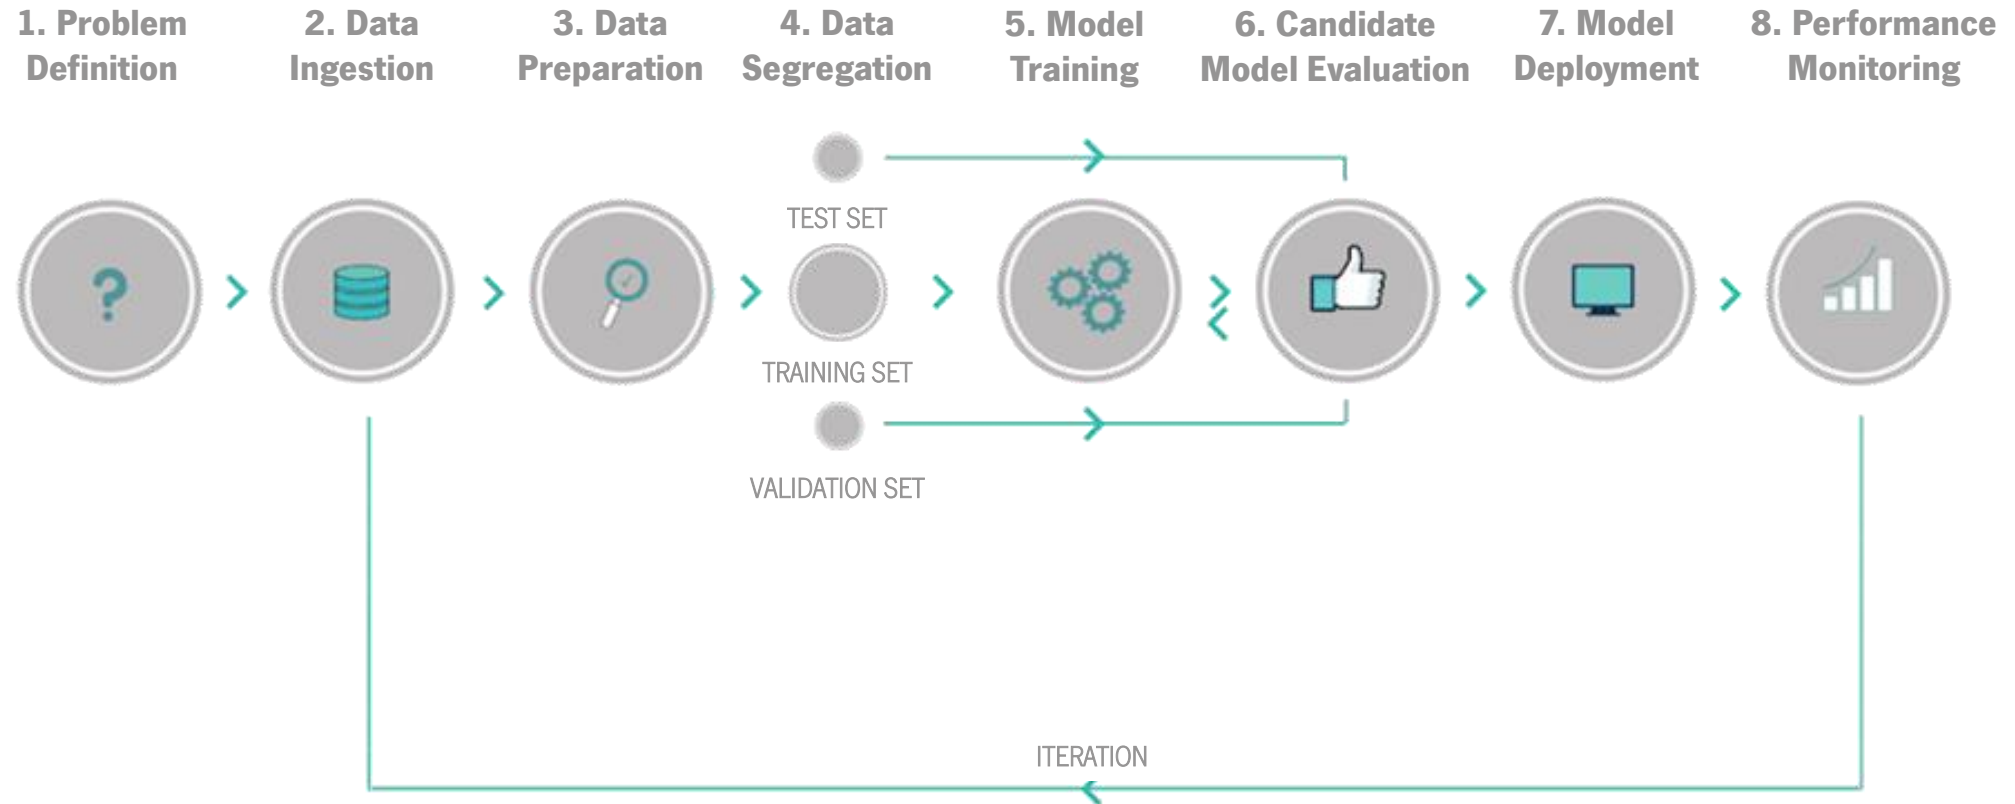
Bruno Fernandes, Dalila Alves, Filipa Ferraz, Victor Alves

Part IX

Contents

2

- Unsupervised Learning
 - K-means Clustering
 - K-medoids Clustering
 - DBSCAN Clustering
- Hands On





Unsupervised Learning

Unsupervised Learning

5

Unsupervised learning means that there is no outcome to be predicted, and the algorithm just tries to find patterns on the data. Using **clustering**, it tries to group (cluster) the data based on their similarity.

For this example, we create the data creating *isotropic Gaussian blobs* from [sklearn.datasets](https://scikit-learn.org/stable/datasets/index.html).

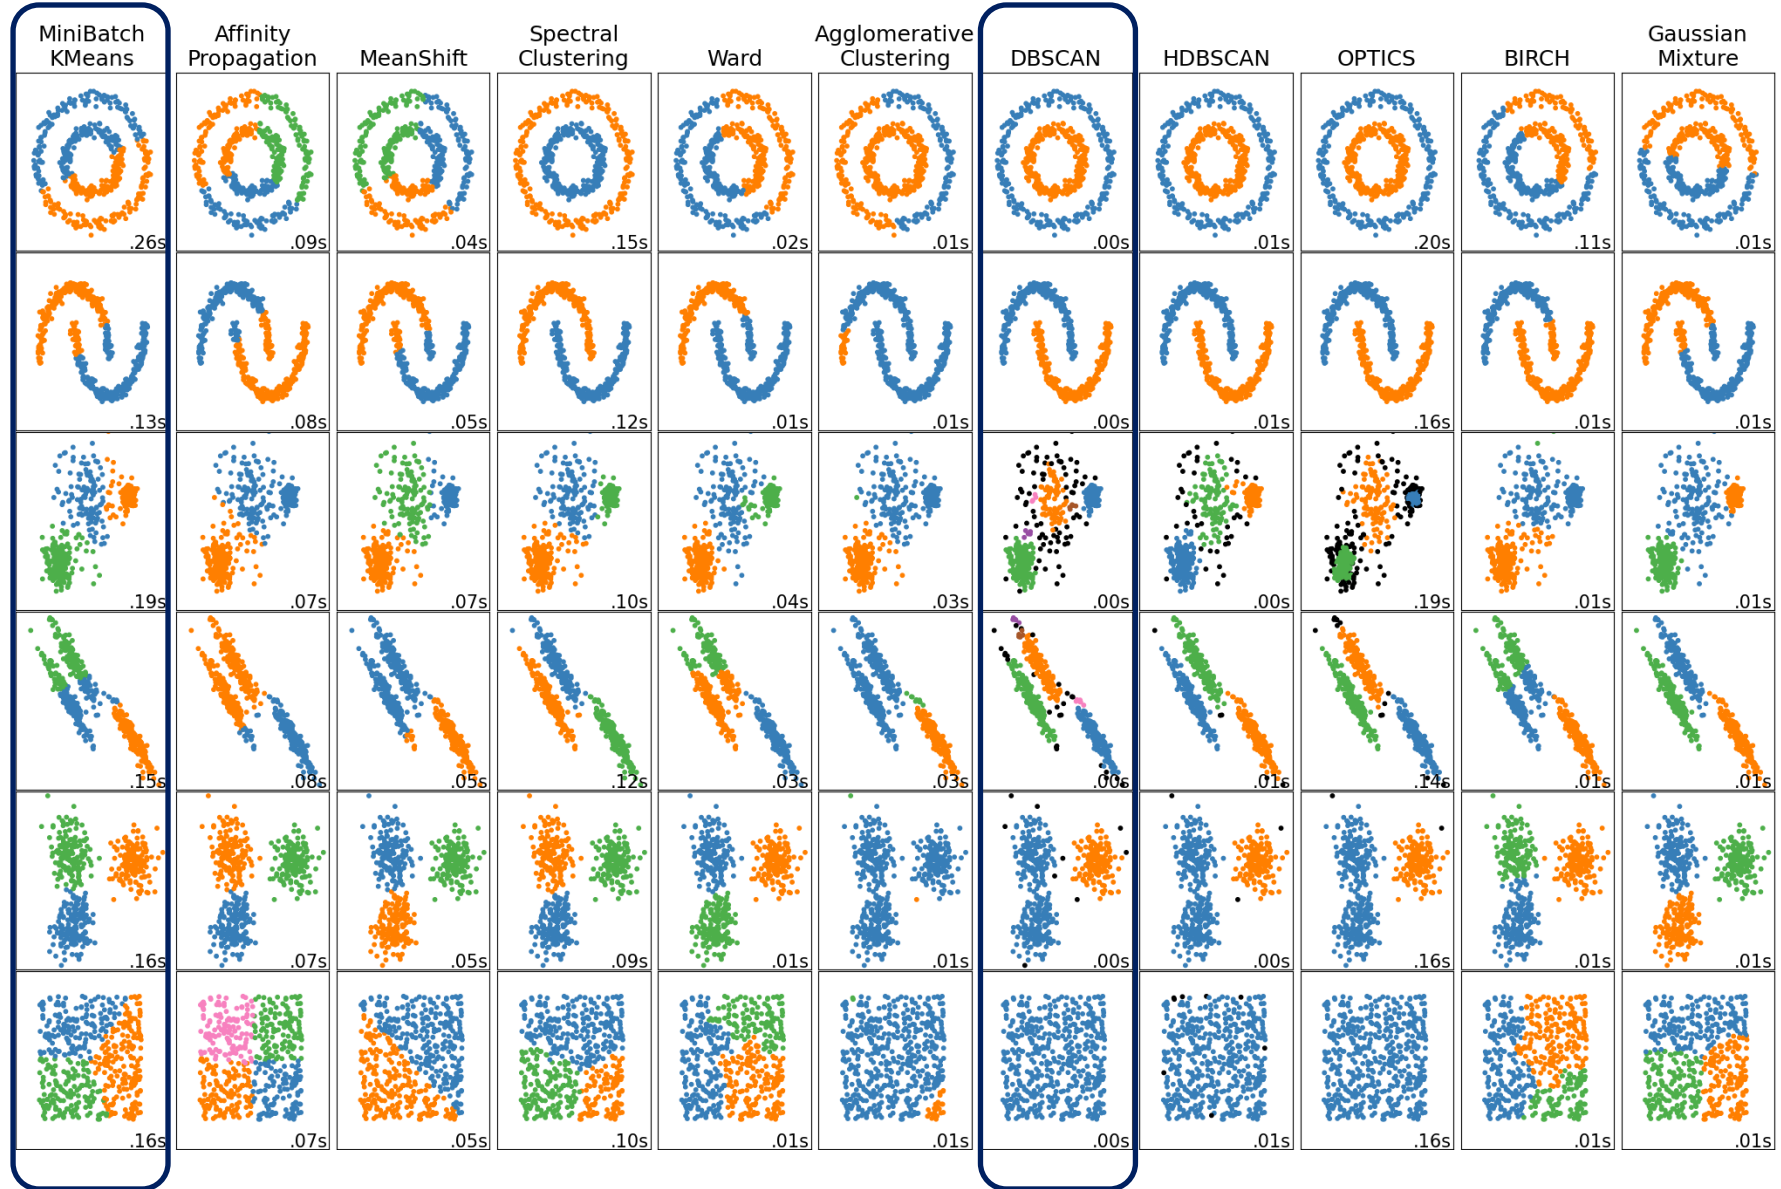
Note: all clustering algorithms require data preprocessing(e.g. dimensionality reduction) and standardization.

You may need to install [scikit-learn-extra](https://github.com/scikit-learn-contrib/scikit-learn-extra).

```
pip install https://github.com/scikit-learn-contrib/scikit-learn-extra/archive/master.zip
```

Unsupervised Learning - Clustering

6




Unsupervised Learning – K-Means and K-Medoids


7

1st iteration

Kmedoids Cluster




Kmeans Cluster




2nd iteration

Kmedoids Cluster



Kmeans Cluster



Unsupervised Learning – K-Means and K-Medoids

8

Both **K-Means** and **K-Medoids** algorithms follow the **partitioning method**. So they are:

- breaking the dataset up into k groups;
- trying to minimize the distance between points of the same cluster and a particular point which is the center of that cluster.

The **K-Means** algorithm chooses *centroids*, the geometric center of a cluster using the mean to all cluster points.

On the other hand, the **K-Medoids** algorithm chooses *medoids*, points as centers that belong to the dataset and computes dissimilarities between data points.

The most common implementation of the K-Medoids clustering algorithm is the **Partitioning Around Medoids (PAM)** algorithm. The PAM algorithm uses a *greedy search*, which may not find the global optimum solution.

Unsupervised Learning – K-Means and K-Medoids

9

Medoids are more robust to outliers than *centroids*, but they need more computation for high dimensional data.

Strengths:



- simple and intuitive;
- scales to large datasets;
- as a result, we also have centroids that can be used as standard cluster representatives

Weaknesses:

- knowledge about the number of clusters is necessary and must be specified as a parameter;
- does not cope well with a very large number of features;
- separates only convex and homogeneous clusters well;
- can result in poor local solutions, so it needs to be run several times.


Unsupervised Learning – K-Means and K-Medoids

10



Unsupervised Learning – K-Means and K-Medoids

11




The Choice of Centriod in K-Means is random and what if there is presence of Outlier?

K-Medoids solves the issue as it K-Medoids selected the precise centriod among all data points of the corresponding clusters


Unsupervised Learning – K-Means and K-Medoids

12



Unsupervised Learning – K-Means and K-Medoids

13



Implementing K-Means

14

Creating the dataset

```
from sklearn.datasets import make_blobs
```

```
data = make_blobs(n_samples = 200, n_features = 2, centers = 4, cluster_std = 1.8, random_state = 2022)
```

Defining X and y

```
X = data[0]  
y = data[1]
```


Visualizing the data (first 5 instances)

```
print('X:', X[0:5, :])  
print('Y:', y[0:5])
```

```
X: [[ 5.88508997  2.9021639 ]  
 [ -8.20429992 -11.68670283]  
 [  1.9125188  -2.76746603]  
 [ -9.39601207 -7.2830252 ]  
 [  6.1986976  7.32152342]]  
Y: [2 1 2 1 3]
```

```
plt.scatter(X[:, 0], X[:, 1], c = y, cmap = 'rainbow')  
for i, txt in enumerate(y):  
    if i%5 == 0:  
        plt.annotate(txt, (X[i, 0], X[i, 1]))
```

4 groups



Implementing K-Means

15

```
sklearn.cluster.KMeans(n_clusters=8, *, init='k-means++', n_init='auto', max_iter=300,  
tol=0.0001, verbose=0, random_state=None, copy_x=True, algorithm='lloyd')
```

```
from sklearn.cluster import KMeans
```

Creating the clusters

```
kmeans = KMeans(n_clusters = 4, n_init = 10, random_state = 2022)
```

```
kmeans.fit(X)
```

```
KMeans  
KMeans(n_clusters=4, n_init=10, random_state=2022)
```

Vector with center of the clusters

```
kmeans.cluster_centers_
```

```
array([[ -7.68797564, -8.88054369],  
       [  7.70499062,  2.96975295],  
       [-9.78544862, -0.2509739 ],  
       [  3.60428123, -0.21752545]])
```

```
kmeans.labels_
```


```
array([[1, 0, 3, 0, 1, 2, 0, 2, 3, 3, 1, 1, 3, 1, 2, 1, 1, 1, 2, 0, 3,  
       1, 1, 0, 0, 3, 0, 3, 3, 1, 0, 1, 3, 1, 3, 0, 1, 0, 2, 1, 2, 1, 3,  
       3, 3, 0, 0, 3, 0, 0, 3, 1, 1, 0, 3, 1, 3, 1, 2, 2, 2, 3, 3, 1, 2,  
       0, 2, 2, 2, 3, 2, 3, 0, 0, 2, 2, 0, 0, 2, 1, 1, 2, 3, 0, 0, 1, 3,  
       2, 0, 3, 1, 2, 3, 3, 1, 3, 2, 3, 3, 0, 1, 2, 0, 2, 3, 2, 3, 2, 2,  
       1, 2, 2, 2, 3, 2, 3, 3, 0, 3, 1, 1, 1, 3, 0, 0, 1, 3, 2, 1, 0, 2,  
       2, 0, 0, 0, 2, 2, 1, 0, 0, 0, 1, 3, 0, 1, 0, 2, 3, 2, 3, 3, 1, 1,  
       1, 1, 1, 1, 0, 0, 0, 0, 0, 3, 1, 2, 0, 0, 1, 2, 3, 1, 1, 0, 3, 1,  
       3, 0, 0, 2, 3, 0, 2, 2, 0, 2, 1, 1, 2, 3, 2, 3, 3, 0, 2, 1, 2, 0,  
       2, 2])
```

Implementing K-Means

16

Visualizing the clusters

```
f, (ax1, ax2) = plt.subplots(1, 2, sharey = True, figsize = (10, 6))
ax1.set_title('K Means')
ax1.scatter(X[:, 0], X[:, 1], c = kmeans.labels_, cmap = 'rainbow')
ax2.set_title("Original")
ax2.scatter(X[:, 0], X[:, 1], c = y, cmap = 'rainbow')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i, 0], X[i, 1]))
```



Note that the colors are meaningless when in reference to the two plots.

Implementing K-Means

17

Align K-Means prediction class with real values

```
y_pred = kmeans.predict(X)
```

y_pred

```
array([1, 0, 3, 0, 1, 2, 0, 2, 3, 3, 1, 1, 3, 1, 2, 1, 1, 1, 1, 2, 0, 3,
       1, 1, 0, 0, 3, 0, 3, 3, 1, 0, 1, 3, 1, 3, 0, 1, 0, 2, 1, 2, 1, 3,
       3, 3, 0, 0, 3, 0, 0, 3, 1, 1, 0, 3, 1, 3, 1, 2, 2, 2, 3, 3, 1, 2,
       0, 2, 2, 2, 3, 2, 3, 0, 0, 2, 2, 0, 0, 2, 1, 1, 2, 3, 0, 0, 1, 3,
       2, 0, 3, 1, 2, 3, 3, 1, 3, 2, 3, 3, 0, 1, 2, 0, 2, 3, 2, 3, 2, 2,
       1, 2, 2, 2, 3, 2, 3, 3, 0, 3, 1, 1, 1, 3, 0, 0, 1, 3, 2, 1, 0, 2,
       2, 0, 0, 0, 2, 2, 1, 0, 0, 0, 1, 3, 0, 1, 0, 2, 3, 2, 3, 3, 1, 1,
       1, 1, 1, 1, 0, 0, 0, 0, 0, 3, 1, 2, 0, 0, 1, 2, 3, 1, 1, 0, 3, 1,
       3, 0, 0, 2, 3, 0, 2, 2, 0, 2, 1, 1, 2, 3, 2, 3, 3, 0, 2, 1, 2, 0,
       2, 2])
```

y

```
array([2, 1, 2, 1, 3, 0, 1, 0, 2, 2, 3, 3, 2, 3, 0, 3, 3, 3, 3, 0, 1, 2,
       3, 3, 1, 1, 3, 1, 2, 2, 3, 1, 3, 3, 3, 2, 1, 3, 1, 0, 3, 0, 3, 2,
       2, 2, 1, 1, 2, 1, 1, 2, 3, 3, 1, 2, 3, 2, 3, 0, 0, 0, 2, 2, 3, 0,
       1, 0, 0, 0, 2, 0, 2, 1, 1, 0, 0, 1, 1, 0, 2, 3, 0, 2, 1, 1, 3, 2,
       0, 1, 2, 3, 0, 2, 2, 3, 2, 0, 2, 2, 1, 3, 0, 1, 0, 2, 0, 2, 0, 0,
       3, 0, 0, 0, 2, 0, 2, 2, 1, 2, 3, 3, 3, 2, 1, 1, 3, 2, 0, 3, 1, 0,
       0, 1, 1, 1, 0, 0, 3, 1, 1, 1, 3, 2, 1, 3, 1, 0, 2, 0, 2, 2, 3, 3,
       2, 3, 3, 3, 0, 1, 1, 1, 1, 2, 3, 0, 1, 1, 3, 0, 2, 3, 3, 1, 2, 3,
       2, 1, 1, 0, 2, 1, 0, 0, 1, 0, 3, 3, 0, 2, 0, 3, 2, 1, 0, 2, 0, 1,
       0, 0])
```

```
y_pred = np.where(y_pred==0, 10, y_pred)
y_pred = np.where(y_pred==2, 0, y_pred)
y_pred = np.where(y_pred==3, 2, y_pred)
y_pred = np.where(y_pred==1, 3, y_pred)
y_pred = np.where(y_pred==10, 1, y_pred)
```

y_pred


```
array([3, 1, 2, 1, 3, 0, 1, 0, 2, 2, 3, 3, 2, 3, 0, 3, 3, 3, 3, 0, 1, 2,
       3, 3, 1, 1, 2, 1, 2, 2, 3, 1, 3, 2, 3, 2, 1, 3, 1, 0, 3, 0, 3, 2,
       2, 2, 1, 1, 2, 1, 1, 2, 3, 3, 1, 2, 3, 2, 3, 0, 0, 0, 2, 2, 3, 0,
       1, 0, 0, 0, 2, 0, 2, 1, 1, 0, 0, 1, 1, 0, 3, 3, 0, 2, 1, 1, 3, 2,
       0, 1, 2, 3, 0, 2, 2, 3, 2, 0, 2, 2, 1, 3, 0, 1, 0, 2, 0, 2, 0, 0,
       3, 0, 0, 0, 2, 0, 2, 2, 1, 2, 3, 3, 3, 2, 1, 1, 3, 2, 0, 3, 1, 0,
       0, 1, 1, 1, 0, 0, 3, 1, 1, 1, 3, 2, 1, 3, 1, 0, 2, 0, 2, 2, 3, 3,
       3, 3, 3, 3, 1, 1, 1, 1, 1, 2, 3, 0, 1, 1, 3, 0, 2, 3, 3, 1, 2, 3,
       2, 1, 1, 0, 2, 1, 0, 0, 1, 0, 3, 3, 0, 2, 0, 2, 2, 1, 0, 3, 0, 1,
       0, 0])
```

Implementing K-Means

18

Redo the visualizations

```
f, (ax1, ax2) = plt.subplots(1, 2, sharey = True, figsize = (8, 5))
ax1.set_title('Original')
ax1.scatter(X[:,0], X[:,1], c=y, cmap='rainbow')
ax2.set_title("Prediction")
ax2.scatter(X[:,0], X[:,1], c=y_pred, cmap='rainbow')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i,0], X[i,1]))
plt.savefig("KMeans_pred.png")
```



Implementing K-Means

19

Evaluate the model

```
print(confusion_matrix(y, y_pred))
```

```
[[49  1  0  0]
 [ 0 50  0  0]
 [ 0  0 46  4]
 [ 0  0  3 47]]
```

```
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	50
1	0.98	1.00	0.99	50
2	0.94	0.92	0.93	50
3	0.92	0.94	0.93	50
accuracy			0.96	200
macro avg	0.96	0.96	0.96	200
weighted avg	0.96	0.96	0.96	200

If you want to classify new points, it is best to train a classifier on your clustering result.

Implementing K-Medoids

20

Let's use the dataset created earlier.

```
sklearn_extra.cluster.KMedoids(n_clusters=8, metric='euclidean', method='alternate',  
init='heuristic', max_iter=300, random_state=None)
```

```
from sklearn_extra.cluster import KMedoids
```

Creating the clusters

```
kmedoids = KMedoids(n_clusters = 4, random_state = 2022)  
kmedoids.fit(X)
```

▼ KMedoids ⓘ
KMedoids(n_clusters=4, random_state=2022)

Vector with center of the clusters

```
kmedoids.cluster_centers_  
  
array([[ -7.62795904, -8.8354951 ],  
       [-9.87312876,  0.13931247],  
       [ 3.65994783, -0.13260646],  
       [ 7.66163195,  2.9333056 ]])
```

```
kmedoids.labels_
```


```
array([3, 0, 2, 0, 3, 1, 0, 1, 2, 2, 3, 3, 2, 3, 1, 3, 3, 3, 3, 1, 0, 2,  
       3, 3, 0, 0, 2, 0, 2, 2, 3, 0, 3, 2, 3, 2, 0, 3, 0, 1, 3, 1, 3, 2,  
       2, 2, 0, 0, 2, 0, 0, 2, 3, 3, 0, 2, 3, 2, 3, 1, 1, 1, 2, 2, 3, 1,  
       0, 1, 1, 1, 2, 1, 2, 0, 0, 1, 1, 0, 0, 1, 3, 3, 1, 2, 0, 0, 3, 2,  
       1, 0, 2, 3, 1, 2, 2, 3, 2, 1, 2, 2, 0, 3, 1, 0, 1, 2, 1, 2, 1, 1,  
       3, 1, 1, 1, 2, 1, 2, 2, 0, 2, 3, 3, 3, 2, 0, 0, 3, 2, 1, 3, 0, 1,  
       1, 0, 0, 0, 1, 1, 3, 0, 0, 0, 3, 2, 0, 3, 0, 1, 2, 1, 2, 2, 3, 3,  
       3, 3, 3, 3, 0, 0, 0, 0, 0, 2, 3, 1, 0, 0, 3, 1, 2, 3, 3, 0, 2, 3,  
       2, 0, 0, 1, 2, 0, 1, 1, 0, 1, 3, 3, 1, 2, 1, 2, 2, 0, 1, 3, 1, 0,  
       1, 1], dtype=int64)
```

Implementing K-Medoids

21

Visualizing the clusters

```
f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(8,5))
ax1.set_title('K Medoids')
ax1.scatter(X[:, 0], X[:, 1], c = kmedoids.labels_, cmap = 'rainbow')
ax2.set_title("Original")
ax2.scatter(X[:, 0], X[:, 1], c = y, cmap = 'rainbow')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i, 0], X[i, 1]))
plt.savefig("KMedoids.png")
```



Implementing K-Medoids

22

Align K-Medoids prediction class with real values

```
y_pred = kmedoids.predict(X)
```

```
y_pred
```

```
array([3, 0, 2, 0, 3, 1, 0, 1, 2, 2, 3, 3, 2, 3, 1, 3, 3, 3, 3, 1, 0, 2,
       3, 3, 0, 0, 2, 0, 2, 2, 3, 0, 3, 2, 3, 2, 0, 3, 0, 1, 3, 1, 3, 2,
       2, 2, 0, 0, 2, 0, 0, 2, 3, 3, 0, 2, 3, 2, 3, 1, 1, 1, 2, 2, 3, 1,
       0, 1, 1, 1, 2, 1, 2, 0, 0, 1, 1, 0, 0, 1, 3, 3, 1, 2, 0, 0, 3, 2,
       1, 0, 2, 3, 1, 2, 2, 3, 2, 1, 2, 2, 0, 3, 1, 0, 1, 2, 1, 2, 1, 1,
       3, 1, 1, 1, 2, 1, 2, 2, 0, 2, 3, 3, 3, 2, 0, 0, 3, 2, 1, 3, 0, 1,
       1, 0, 0, 0, 1, 1, 3, 0, 0, 0, 3, 2, 0, 3, 0, 1, 2, 1, 2, 2, 3, 3,
       3, 3, 3, 3, 0, 0, 0, 0, 0, 2, 3, 1, 0, 0, 3, 1, 2, 3, 3, 0, 2, 3,
       2, 0, 0, 1, 2, 0, 1, 1, 0, 1, 3, 3, 1, 2, 1, 2, 2, 0, 1, 3, 1, 0,
       1, 1], dtype=int64)
```

```
y
```

```
array([2, 1, 2, 1, 3, 0, 1, 0, 2, 2, 3, 3, 2, 3, 0, 3, 3, 3, 3, 0, 1, 2,
       3, 3, 1, 1, 3, 1, 2, 2, 3, 1, 3, 3, 3, 2, 1, 3, 1, 0, 3, 0, 3, 2,
       2, 2, 1, 1, 2, 1, 1, 2, 3, 3, 1, 2, 3, 2, 3, 0, 0, 0, 2, 2, 3, 0,
       1, 0, 0, 0, 2, 0, 2, 1, 1, 0, 0, 1, 1, 0, 2, 3, 0, 2, 1, 1, 3, 2,
       0, 1, 2, 3, 0, 2, 2, 3, 2, 0, 2, 2, 1, 3, 0, 1, 0, 2, 0, 2, 0, 0,
       3, 0, 0, 0, 2, 0, 2, 2, 1, 2, 3, 3, 3, 2, 1, 1, 3, 2, 0, 3, 1, 0,
       0, 1, 1, 1, 0, 0, 3, 1, 1, 1, 3, 2, 1, 3, 1, 0, 2, 0, 2, 2, 3, 3,
       2, 3, 3, 3, 0, 1, 1, 1, 1, 2, 3, 0, 1, 1, 3, 0, 2, 3, 3, 1, 2, 3,
       2, 1, 1, 0, 2, 1, 0, 0, 1, 0, 3, 3, 0, 2, 0, 3, 2, 1, 0, 2, 0, 1,
       0, 0])
```

```
y_pred = np.where(y_pred==1, 10, y_pred)
y_pred = np.where(y_pred==0, 1, y_pred)
y_pred = np.where(y_pred==10, 0, y_pred)
```

```
y_pred
```


```
array([3, 1, 2, 1, 3, 0, 1, 0, 2, 2, 3, 3, 2, 3, 0, 3, 3, 3, 3, 0, 1, 2,
       3, 3, 1, 1, 2, 1, 2, 2, 3, 1, 3, 2, 3, 2, 1, 3, 1, 0, 3, 0, 3, 2,
       2, 2, 1, 1, 2, 1, 1, 2, 3, 3, 1, 2, 3, 2, 3, 0, 0, 0, 2, 2, 3, 0,
       1, 0, 0, 0, 2, 0, 2, 1, 1, 0, 0, 1, 1, 0, 3, 3, 0, 2, 1, 1, 3, 2,
       0, 1, 2, 3, 0, 2, 2, 3, 2, 0, 2, 2, 1, 3, 0, 1, 0, 2, 0, 2, 0, 0,
       3, 0, 0, 0, 2, 0, 2, 2, 1, 2, 3, 3, 3, 2, 1, 1, 3, 2, 0, 3, 1, 0,
       0, 1, 1, 1, 0, 0, 3, 1, 1, 1, 3, 2, 1, 3, 1, 0, 2, 0, 2, 2, 3, 3,
       3, 3, 3, 3, 1, 1, 1, 1, 1, 2, 3, 0, 1, 1, 3, 0, 2, 3, 3, 1, 2, 3,
       2, 1, 1, 0, 2, 1, 0, 0, 1, 0, 3, 3, 0, 2, 0, 2, 2, 1, 0, 3, 0, 1,
       0, 0], dtype=int64)
```

Implementing K-Medoids

23

Redo the visualizations

```
f, (ax1, ax2) = plt.subplots(1, 2, sharey = True, figsize = (8, 5))
ax1.set_title('Original')
ax1.scatter(X[:, 0], X[:, 1], c = y, cmap = 'rainbow')
ax2.set_title("Prediction")
ax2.scatter(X[:, 0], X[:, 1], c = y_pred, cmap = 'rainbow')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i, 0], X[i, 1]))
plt.savefig("KMedoid_pred.png")
```



Implementing K-Medoids

24

Evaluate the model

```
print(confusion_matrix(y, y_pred))
```

```
[[49  1  0  0]
 [ 0 50  0  0]
 [ 0  0 46  4]
 [ 0  0  3 47]]
```

```
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	50
1	0.98	1.00	0.99	50
2	0.94	0.92	0.93	50
3	0.92	0.94	0.93	50
accuracy			0.96	200
macro avg	0.96	0.96	0.96	200
weighted avg	0.96	0.96	0.96	200

K-Means vs. K-Medoids

25

K-Means

```
print(confusion_matrix(y, y_pred))
```

```
[[49  1  0  0]
 [ 0 50  0  0]
 [ 0  0 46  4]
 [ 0  0  3 47]]
```

```
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	50
1	0.98	1.00	0.99	50
2	0.94	0.92	0.93	50
3	0.92	0.94	0.93	50
accuracy			0.96	200
macro avg	0.96	0.96	0.96	200
weighted avg	0.96	0.96	0.96	200

K-Medoids

```
print(confusion_matrix(y, y_pred))
```

```
[[49  1  0  0]
 [ 0 50  0  0]
 [ 0  0 46  4]
 [ 0  0  3 47]]
```

```
print(classification_report(y, y_pred))
```

	precision	recall	f1-score	support
0	1.00	0.98	0.99	50
1	0.98	1.00	0.99	50
2	0.94	0.92	0.93	50
3	0.92	0.94	0.93	50
accuracy			0.96	200
macro avg	0.96	0.96	0.96	200
weighted avg	0.96	0.96	0.96	200

Can we compare these two models?

Unsupervised Learning - DBSCAN

26

The **DBSCAN** (Density-Based Spatial Clustering of Applications with Noise) algorithm finds core samples in high-density regions and expands clusters for them.

High-density regions, where data points are located close to each other, are separated by **low-density regions**, where the data points are located far from each other.


The idea of a **core sample** means a sample located in an area of high-density. Data point A is considered a core sample if at least minimum number of points required to form a dense region (usually including A) are located within ϵ distance from A.




Unsupervised Learning - DBSCAN

27

Influence of the size of ϵ



Influence of the density of the region



Unsupervised Learning - DBSCAN

28

Strengths:

- knowledge about the number of clusters is not necessary;
- also solves the anomaly detection task.

Weaknesses:

- need to select and tune the density parameter (ϵ);
- does not cope well with sparse data.

Implementing DBSCAN

29

Creating the dataset

```
X, y = make_blobs(n_samples=750, cluster_std=0.4, random_state=0)
```

Shape of X and y and visualizing the data

X.shape

y.shape

(750, 2)

(750,)

```
print('X:', X[0:5])
```


```
print('y:', y[0:5])
```

```
X: [[ 2.36434546  0.23302434]
 [ 0.92311785  4.18467098]
 [ 1.64221028  0.72296432]
 [ 1.97590796  0.93534058]
 [-1.68752703  2.73049184]]
```

```
y: [1 0 1 1 2]
```

```
plt.figure(figsize=(8, 6))
plt.scatter(X[:, 0], X[:, 1], c = y, cmap = 'viridis')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i, 0], X[i, 1]))
plt.title('Blobs Dataset')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

3 groups



Implementing DBSCAN

30

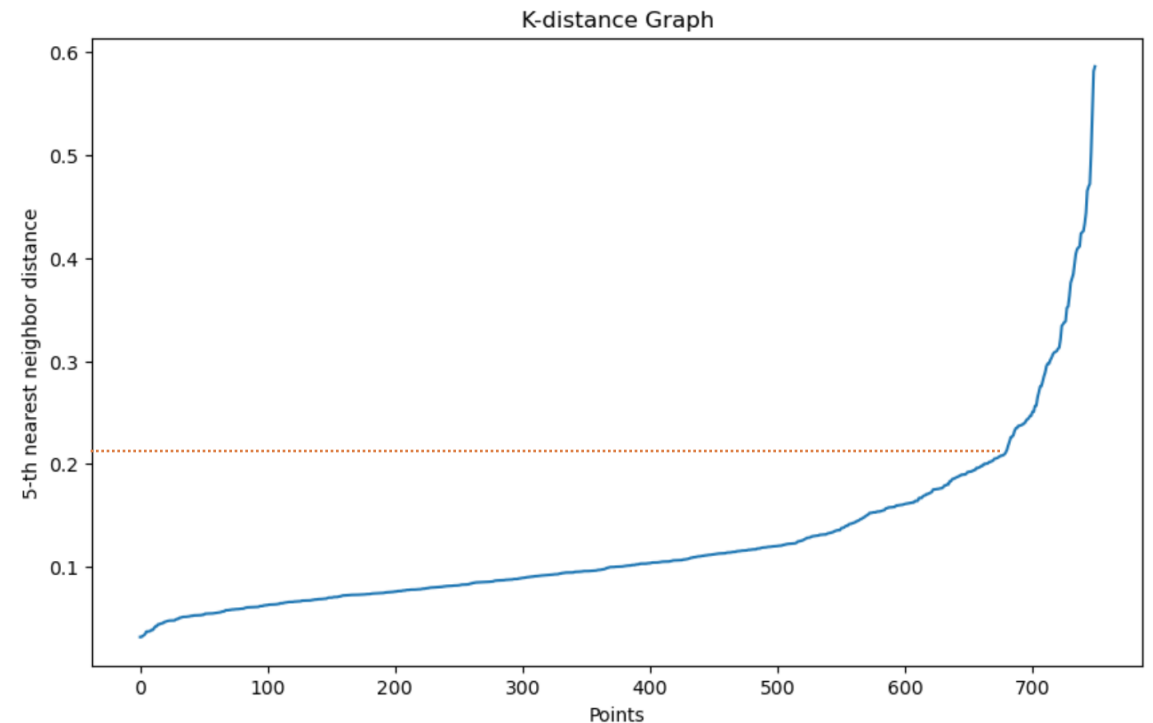
Determining the epsilon parameter

We use the k -distance graph method to help choose an appropriate ε value:

1. We define a function `plot_k_distance_graph` that calculates the distance to the k^{th} nearest neighbor for each point
2. The distances are sorted and plotted.
3. We look for an "elbow" in the resulting graph to choose ε .

```
from sklearn.neighbors import NearestNeighbors

def plot_k_distance_graph(X, k):
    neigh = NearestNeighbors(n_neighbors=k)
    neigh.fit(X)
    distances, _ = neigh.kneighbors(X)
    distances = np.sort(distances[:, k-1])
    plt.figure(figsize=(10, 6))
    plt.plot(distances)
    plt.xlabel('Points')
    plt.ylabel(f'{k}-th nearest neighbor distance')
    plt.title('K-distance Graph')
    plt.show()
plot_k_distance_graph(X, k=5)
```



Implementing DBSCAN

31

```
sklearn.cluster.DBSCAN(eps=0.5, *, min_samples=5, metric='euclidean', metric_params=None,  
algorithm='auto', leaf_size=30, p=None, n_jobs=None)
```

```
from sklearn.cluster import DBSCAN
```

Creating the clusters

```
db = DBSCAN(eps=0.2, min_samples=10)
```

```
db.fit(X)
```

```
▼ DBSCAN ⓘ ?  
DBSCAN(eps=0.2, min_samples=10)
```

Confirm the labels shape

```
labels = db.labels_  
labels.shape
```

```
(750,)
```

Implementing DBSCAN

32

Interpreting the results

Indices of core samples

```
db.core_sample_indices_  
  
array([ 1,  2,  3,  4,  6,  8, 10, 12, 14, 15, 16, 18, 20,  
       ...,  
       736, 738, 739, 740, 741, 742, 743, 745, 746, 748, 749], dtype=int64)
```

Components

```
db.components_  
  
array([[0.92311785, 4.18467098],  
       [1.64221028, 0.72296432],  
       [1.97590796, 0.93534058],  
       ...,  
       [1.04724054, 4.14307495],  
       [2.16897939, 1.59473117],  
       [0.73774446, 4.28276041]])
```

Number of features

```
db.n_features_in_  
  
2
```

Number of clusters and noise points in *labels*

```
n_clusters_ = len(set(labels)) - (1 if -1 in labels else 0)  
n_noise_ = list(labels).count(-1)
```


```
print("Estimated number of clusters: %d" % n_clusters_)  
print("Estimated number of noise points: %d" % n_noise_)
```

```
Estimated number of clusters: 3  
Estimated number of noise points: 104
```


Implementing DBSCAN

33

Visualizing the resultant clusters




```
unique_labels = set(labels)
core_samples_mask = np.zeros_like(labels, dtype=bool)
core_samples_mask[db.core_sample_indices_] = True

plt.figure(figsize=(8, 6))
colors = [plt.cm.Spectral(each) for each in np.linspace(0, 1, len(unique_labels))]
for k, col in zip(unique_labels, colors):
    if k == -1:
        col = [0, 0, 0, 1]
    class_member_mask = labels == k
    xy = X[class_member_mask & core_samples_mask]
    plt.plot(
        xy[:, 0],
        xy[:, 1],
        "o",
        markerfacecolor=tuple(col),
        markeredgecolor="k",
        markersize=14,)
    xy = X[class_member_mask & ~core_samples_mask]
    plt.plot(
        xy[:, 0],
        xy[:, 1],
        "o",
        markerfacecolor=tuple(col),
        markeredgecolor="k",
        markersize=6,)
plt.title(f"Estimated number of clusters: {n_clusters}")
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```

Implementing DBSCAN

34

```
plt.figure(figsize=(10, 6))
scatter = plt.scatter(X[:, 0], X[:, 1], c=pred, cmap='viridis')
plt.colorbar(scatter)
plt.title('DBSCAN Clustering Results')
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')
plt.show()
```



Implementing DBSCAN

35

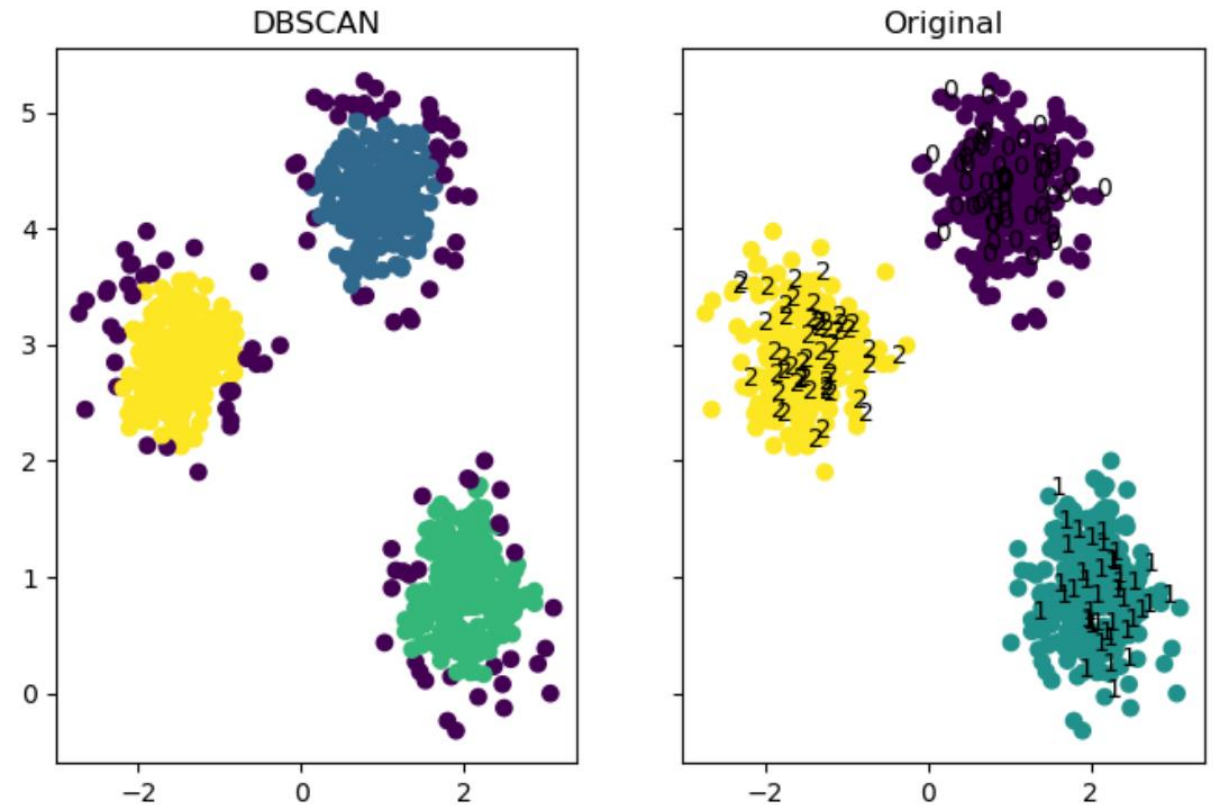
Evaluating the model

Obtain the predictions

```
pred = db.fit_predict(X)
```

Visualize the original vs. DBSCAN clusters

```
f, (ax1, ax2) = plt.subplots(1, 2, sharey=True, figsize=(8,5))
ax1.set_title('DBSCAN')
ax1.scatter(X[:, 0], X[:, 1], c = db.labels_, cmap = 'viridis')
ax2.set_title("Original")
ax2.scatter(X[:, 0], X[:, 1], c = y, cmap = 'viridis')
for i, txt in enumerate(y):
    if i%5 == 0:
        plt.annotate(txt, (X[i, 0], X[i, 1]))
plt.savefig("DBSCAN.png")
plt.show()
```



Implementing DBSCAN

36

Because Blobs provides the true labels, we can analyze through the following metrics: homogeneity, completeness, V-measure, Adjusted Rand Index, Adjusted Mutual Information and Silhouette Coefficient

```
print(f"Homogeneity: {metrics.homogeneity_score(y, labels):.3f}")
print(f"Completeness: {metrics.completeness_score(y, labels):.3f}")
print(f"V-measure: {metrics.v_measure_score(y, labels):.3f}")
print(f"Adjusted Rand Index: {metrics.adjusted_rand_score(y, labels):.3f}")
print(
    "Adjusted Mutual Information:"
    f" {metrics.adjusted_mutual_info_score(y, labels):.3f}"
)
print(f"Silhouette Coefficient: {metrics.silhouette_score(X, labels):.3f}")
```

```
Homogeneity: 0.862
Completeness: 0.702
V-measure: 0.774
Adjusted Rand Index: 0.781
Adjusted Mutual Information: 0.773
Silhouette Coefficient: 0.586
```

```
print(confusion_matrix(y, pred))
```

```
[[ 0  0  0  0]
 [39 211  0  0]
 [31  0 219  0]
 [34  0  0 216]]
```

```
print(classification_report(y, pred))
```

	precision	recall	f1-score	support
-1	0.00	0.00	0.00	0
0	1.00	0.84	0.92	250
1	1.00	0.88	0.93	250
2	1.00	0.86	0.93	250
accuracy			0.86	750
macro avg	0.75	0.65	0.69	750
weighted avg	1.00	0.86	0.93	750

What do these metrics mean?

Implementing DBSCAN

37


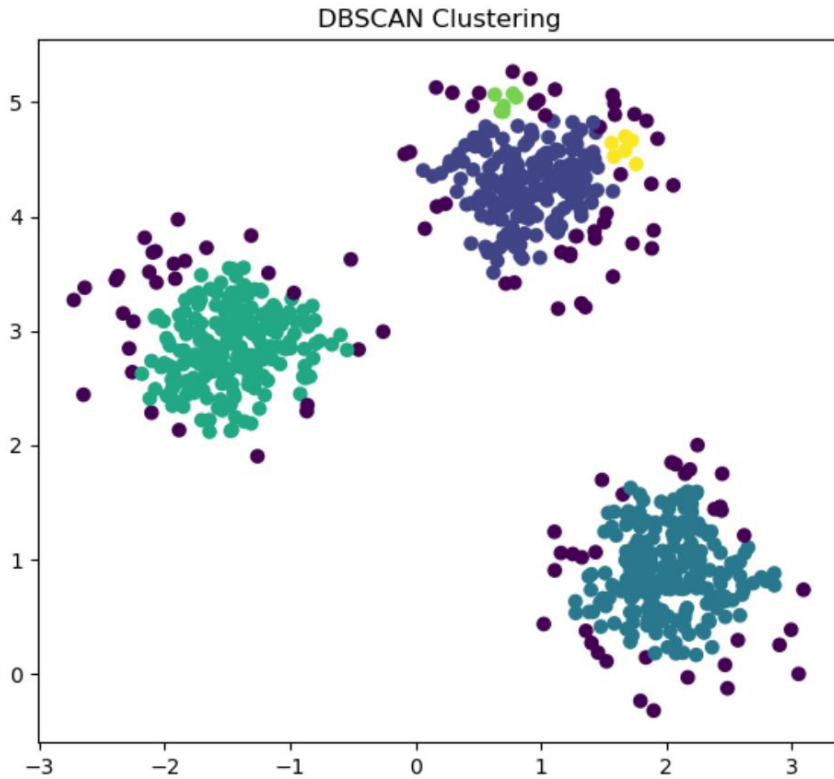
Comparing DBSCAN and K-Means

```
dbscan = DBSCAN(eps=0.15, min_samples=5)
dbscan_labels = dbscan.fit_predict(X)
kmeans = KMeans(n_clusters=3, random_state=42)
kmeans_labels = kmeans.fit_predict(X)
fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(15, 6))
ax1.scatter(X[:, 0], X[:, 1], c=dbscan_labels, cmap='viridis')
ax1.set_title('DBSCAN Clustering')
ax2.scatter(X[:, 0], X[:, 1], c=kmeans_labels, cmap='viridis')
ax2.set_title('K-Means Clustering')
plt.show()
```

Implementing DBSCAN

38

Comparing DBSCAN and K-Means



What conclusions can be obtained?



Hands On